

# General

## Q: What is AWS Lambda?

AWS Lambda lets you run code without provisioning or managing servers. You pay only for the compute time you consume - there is no charge when your code is not running. With Lambda, you can run code for virtually any type of application or backend service - all with zero administration. Just upload your code and Lambda takes care of everything required to run and scale your code with high availability. You can set up your code to automatically trigger from other AWS services or call it directly from any web or mobile app.

## Q: What is serverless computing?

[Serverless computing](#) allows you to build and run applications and services without thinking about servers. With serverless computing, your application still runs on servers, but all the server management is done by AWS. At the core of serverless computing is AWS Lambda, which lets you run your code without provisioning or managing servers.

## Q: What events can trigger an AWS Lambda function?

Please see our [documentation](#) for a complete list of event sources.

## Q: When should I use AWS Lambda versus Amazon EC2?

Amazon Web Services offers a set of compute services to meet a range of needs.

[Amazon EC2](#) offers flexibility, with a wide range of instance types and the option to customize the operating system, network and security settings, and the entire software stack, allowing you to easily move existing applications to the cloud. With Amazon EC2 you are responsible for provisioning capacity, monitoring fleet health and performance, and designing for fault tolerance and scalability. [AWS Elastic Beanstalk](#) offers an easy-to-use service for deploying and scaling web applications in which you retain ownership and full control over the underlying EC2 instances. [Amazon EC2 Container Service](#) is a scalable management service that supports Docker containers and allows you to easily run distributed applications on a managed cluster of Amazon EC2 instances.

AWS Lambda makes it easy to execute code in response to events, such as changes to Amazon S3 buckets, updates to an Amazon DynamoDB table, or custom events generated by your applications or devices. With Lambda you do not have to provision your own instances; Lambda performs all the operational and administrative activities on your behalf, including capacity provisioning, monitoring fleet health, applying security patches to the underlying compute resources, deploying your code, running a web service front end, and monitoring and logging

your code. AWS Lambda provides easy scaling and high availability to your code without additional effort on your part.

#### **Q: What kind of code can run on AWS Lambda?**

AWS Lambda offers an easy way to accomplish many activities in the cloud. For example, you can use AWS Lambda to build mobile back-ends that retrieve and transform data from Amazon DynamoDB, handlers that compress or transform objects as they are uploaded to Amazon S3, auditing and reporting of API calls made to any Amazon Web Service, and server-less processing of streaming data using Amazon Kinesis.

#### **Q: What languages does AWS Lambda support?**

AWS Lambda natively supports Java, Go, PowerShell, Node.js, C#, Python, and Ruby code, and provides a Runtime API which allows you to use any additional programming languages to author your functions. Please read our documentation on using [Node.js](#), [Python](#), [Java](#), [Ruby](#), [C#](#), [Go](#) and [PowerShell](#).

#### **Q: Can I access the infrastructure that AWS Lambda runs on?**

No. AWS Lambda operates the compute infrastructure on your behalf, allowing it to perform health checks, apply security patches, and do other routine maintenance.

#### **Q: How does AWS Lambda isolate my code?**

Each AWS Lambda function runs in its own isolated environment, with its own resources and file system view. AWS Lambda uses the same techniques as Amazon EC2 to provide security and separation at the infrastructure and execution levels.

#### **Q: How does AWS Lambda secure my code?**

AWS Lambda stores code in Amazon S3 and encrypts it at rest. AWS Lambda performs additional integrity checks while your code is in use.

#### **Q: What AWS regions are available for AWS Lambda?**

Please refer to the [AWS Global Infrastructure Region Table](#).

## **AWS Lambda functions**

#### **Q: What is an AWS Lambda function?**

The code you run on AWS Lambda is uploaded as a “Lambda function”. Each function has associated configuration information, such as its name, description, entry point, and resource requirements. The code must be written in a “stateless” style i.e. it should assume there is no affinity to the underlying compute infrastructure. Local file system access, child processes, and similar artifacts may not extend beyond the lifetime of the request, and any persistent state should be stored in Amazon S3, Amazon DynamoDB, or another Internet-available storage service. Lambda functions can include libraries, even native ones.

#### **Q: Will AWS Lambda reuse function instances?**

To improve performance, AWS Lambda may choose to retain an instance of your function and reuse it to serve a subsequent request, rather than creating a new copy. To learn more about how Lambda reuses function instances, visit our [documentation](#). Your code should not assume that this will always happen.

#### **Q: What if I need scratch space on disk for my AWS Lambda function?**

Each Lambda function receives 500MB of non-persistent disk space in its own /tmp directory.

#### **Q: Why must AWS Lambda functions be stateless?**

Keeping functions stateless enables AWS Lambda to rapidly launch as many copies of the function as needed to scale to the rate of incoming events. While AWS Lambda’s programming model is stateless, your code can access stateful data by calling other web services, such as Amazon S3 or Amazon DynamoDB.

#### **Q: Can I use threads and processes in my AWS Lambda function code?**

Yes. AWS Lambda allows you to use normal language and operating system features, such as creating additional threads and processes. Resources allocated to the Lambda function, including memory, execution time, disk, and network use, must be shared among all the threads/processes it uses. You can launch processes using any language supported by Amazon Linux.

#### **Q: What restrictions apply to AWS Lambda function code?**

Lambda attempts to impose as few restrictions as possible on normal language and operating system activities, but there are a few activities that are disabled: Inbound network connections are blocked by AWS Lambda, and for outbound connections only TCP/IP and UDP/IP sockets are supported, and ptrace (debugging) system calls are blocked. TCP port 25 traffic is also blocked as an anti-spam measure.

#### **Q: How do I create an AWS Lambda function using the Lambda console?**

If you are using Node.js or Python, you can author the code for your function using code editor in the AWS Lambda console which lets you author and test your functions, and view the results of function executions in a robust, IDE-like environment. [Go to the console to get started.](#)

You can also package the code (and any dependent libraries) as a ZIP and upload it using the AWS Lambda console from your local environment or specify an Amazon S3 location where the ZIP file is located. Uploads must be no larger than 50MB (compressed). You can use the AWS Eclipse plugin to author and deploy Lambda functions in Java. You can use the Visual Studio plugin to author and deploy Lambda functions in C#, and Node.js.

## **Q: How do I create an AWS Lambda function using the Lambda CLI?**

You can package the code (and any dependent libraries) as a ZIP and upload it using the AWS CLI from your local environment, or specify an Amazon S3 location where the ZIP file is located. Uploads must be no larger than 50MB (compressed). Visit the [Lambda Getting Started guide](#) to get started.

## **Q: Does AWS Lambda support environment variables?**

Yes. You can easily create and modify environment variables from the AWS Lambda Console, CLI or SDKs. To learn more about environment variables, see the [documentation](#).

## **Q: Can I store sensitive information in environment variables?**

For sensitive information, such as database passwords, we recommend you use client-side encryption using [AWS Key Management Service](#) and store the resulting values as ciphertext in your environment variable. You will need to include logic in your AWS Lambda function code to decrypt these values.

## **Q: How can I manage my AWS Lambda functions?**

You can easily list, delete, update, and monitor your Lambda functions using the dashboard in the AWS Lambda console. You can also use the AWS CLI and AWS SDK to manage your Lambda functions. Visit the [Lambda Developer Guide](#) to learn more.

## **Q: Can I share code across functions?**

Yes, you can package any code (frameworks, SDKs, libraries, and more) as a [Lambda Layer](#) and manage and share them easily across multiple functions.

## **Q: How do I monitor an AWS Lambda function?**

AWS Lambda automatically monitors Lambda functions on your behalf, reporting real-time metrics through Amazon CloudWatch, including total requests, account-level and function-level

concurrency usage, latency, error rates, and throttled requests. You can view statistics for each of your Lambda functions via the Amazon CloudWatch console or through the AWS Lambda console. You can also call third-party monitoring APIs in your Lambda function.

Visit [Troubleshooting CloudWatch metrics](#) to learn more. Standard charges for AWS Lambda apply to use Lambda's built-in metrics.

### **Q: How do I troubleshoot failures in an AWS Lambda function?**

AWS Lambda automatically integrates with Amazon CloudWatch logs, creating a log group for each Lambda function and providing basic application lifecycle event log entries, including logging the resources consumed for each use of that function. You can easily insert additional logging statements into your code. You can also call third-party logging APIs in your Lambda function. Visit [Troubleshooting Lambda functions](#) to learn more. Amazon CloudWatch Logs rates will apply.

### **Q: How do I scale an AWS Lambda function?**

You do not have to scale your Lambda functions – AWS Lambda scales them automatically on your behalf. Every time an event notification is received for your function, AWS Lambda quickly locates free capacity within its compute fleet and runs your code. Since your code is stateless, AWS Lambda can start as many copies of your function as needed without lengthy deployment and configuration delays. There are no fundamental limits to scaling a function. AWS Lambda will dynamically allocate capacity to match the rate of incoming events.

### **Q: How are compute resources assigned to an AWS Lambda function?**

In the AWS Lambda resource model, you choose the amount of memory you want for your function, and are allocated proportional CPU power and other resources. For example, choosing 256MB of memory allocates approximately twice as much CPU power to your Lambda function as requesting 128MB of memory and half as much CPU power as choosing 512MB of memory. To learn more, see our [Function Configuration documentation](#).

You can set your memory from 128MB to 10,240MB.

### **Q: When should I use AWS Lambda functions with more than 3008 MB of memory?**

Customers running memory or compute intensive workloads can now powerup their functions. Larger memory functions help multithreaded applications run faster, making them ideal for data and computationally intensive applications like machine learning, batch and ETL jobs, financial modelling, genomics, HPC, and media processing.

### **Q: How long can an AWS Lambda function execute?**

AWS Lambda functions can be configured to run up to 15 minutes per execution. You can set the timeout to any value between 1 second and 15 minutes.

#### **Q: How will I be charged for using AWS Lambda functions?**

AWS Lambda is priced on a pay per use basis. Please see the [AWS Lambda pricing page](#) for details.

#### **Q: Can I save money on AWS Lambda with a Compute Savings Plan?**

Yes. In addition to saving money on Amazon EC2 and AWS Fargate, you can also use Compute Savings Plans to save money on AWS Lambda. Compute Savings Plans offer up to 17% discount on Duration, Provisioned Concurrency, and Duration (Provisioned Concurrency). Compute Savings Plans do not offer a discount on Requests in your Lambda bill. However, your Compute Savings Plans commitment can apply to Requests at regular rates.

#### **Q: Does AWS Lambda support versioning?**

Yes. By default, each AWS Lambda function has a single, current version of the code. Clients of your Lambda function can call a specific version or get the latest implementation. Please read our documentation on [versioning Lambda functions](#).

#### **Q: How long after uploading my code will my AWS Lambda function be ready to call?**

Deployment times may vary with the size of your code, but AWS Lambda functions are typically ready to call within seconds of upload.

#### **Q: Can I use my own version of a supported library?**

Yes. You can include your own copy of a library (including the AWS SDK) in order to use a different version than the default one provided by AWS Lambda.

## **Using AWS Lambda to process AWS events**

#### **Q: What is an event source?**

An event source is an AWS service or developer-created application that produces events that trigger an AWS Lambda function to run. Some services publish these events to Lambda by invoking the cloud function directly (for example, Amazon S3). Lambda can also poll resources in other services that do not publish events to Lambda. For example, Lambda can pull records from an Amazon Kinesis stream or an Amazon SQS queue and execute a Lambda function for each fetched message.

Many other services, such as AWS CloudTrail, can act as event sources simply by logging to Amazon S3 and using S3 bucket notifications to trigger AWS Lambda functions.

**Q: What event sources can be used with AWS Lambda?**

Please see our [documentation](#) for a complete list of event sources.

**Q: How are events represented in AWS Lambda?**

Events are passed to a Lambda function as an event input parameter. For event sources where events arrive in batches, such as Amazon SQS, Amazon Kinesis, and Amazon DynamoDB Streams, the event parameter may contain multiple events in a single call, based on the batch size you request. To learn more about Amazon S3 event notifications visit [Configuring Notifications for Amazon S3 Events](#). To learn more about Amazon DynamoDB Streams visit the [DynamoDB Stream Developers Guide](#). To learn more about invoking Lambda functions using Amazon SNS, visit the [Amazon SNS Developers Guide](#). For more information on Amazon Cognito events, visit [Amazon Cognito](#). For more information on AWS CloudTrail logs and auditing API calls across AWS services, see [AWS CloudTrail](#).

**Q: How do I make an AWS Lambda function respond to changes in an Amazon S3 bucket?**

From the AWS Lambda console, you can select a function and associate it with notifications from an Amazon S3 bucket. Alternatively, you can use the Amazon S3 console and configure the bucket's notifications to send to your AWS Lambda function. This same functionality is also available through the AWS SDK and CLI.

**Q: How do I make an AWS Lambda function respond to updates in an Amazon DynamoDB table?**

You can trigger a Lambda function on DynamoDB table updates by subscribing your Lambda function to the DynamoDB Stream associated with the table. You can associate a DynamoDB Stream with a Lambda function using the Amazon DynamoDB console, the AWS Lambda console or Lambda's registerEventSource API.

**Q: How do I use an AWS Lambda function to process records in an Amazon Kinesis stream?**

From the AWS Lambda console, you can select a Lambda function and associate it with an Amazon Kinesis stream owned by the same account. This same functionality is also available through the AWS SDK and CLI.

**Q: How does AWS Lambda process data from Amazon Kinesis streams and Amazon DynamoDB Streams?**

The Amazon Kinesis and DynamoDB Streams records sent to your AWS Lambda function are strictly serialized, per shard. This means that if you put two records in the same shard, Lambda guarantees that your Lambda function will be successfully invoked with the first record before it is invoked with the second record. If the invocation for one record times out, is throttled, or encounters any other error, Lambda will retry until it succeeds (or the record reaches its 24-hour expiration) before moving on to the next record. The ordering of records across different shards is not guaranteed, and processing of each shard happens in parallel.

## Q: How should I choose between AWS Lambda and Amazon Kinesis Data Analytics for my analytics needs?

AWS Lambda allows you to perform time-based aggregations (such as count, max, sum, average, etc.) over a short window of up to 15 minutes for your data in Amazon Kinesis or Amazon DynamoDB Streams, over a single logical partition such as a shard. This gives you the option to easily set up simple analytics for your event-based application without adding architectural complexity, as your business and analytics logic can be located in the same function. Lambda allows aggregations over a maximum of a 15-minute tumbling window, based on the event timestamp. [Amazon Kinesis Data Analytics](#) allows you to build more complex analytics applications that support flexible processing choices and robust fault-tolerance with exactly-once processing without duplicates, and analytics that can be performed over an entire data stream across multiple logical partitions. With KDA, you can analyze data over multiple types of aggregation windows (tumbling window, stagger window, sliding window, session window) using either the event time or the processing time.

	AWS Lambda	Amazon KDA
Tumbling Window	Yes	Yes
Stagger Window	No	Yes
Sliding Window	No	Yes
Session Window	No	Yes
Enrichment	No	Yes
Joint input and reference tables	No	Yes
Split input stream	No	Yes
Exactly-once processing	No	Yes
Maximum time window	15 mins	No limit
Aggregation scope	Partition/shard	Stream
Time semantics	Event time	Event time, Processing time

## **Q: How do I use an AWS Lambda function to respond to notifications sent by Amazon Simple Notification Service (SNS)?**

From the AWS Lambda console, you can select a Lambda function and associate it with an Amazon SNS topic. This same functionality is also available through the AWS SDK and CLI.

## **Q: How do I use an AWS Lambda function to respond to emails sent by Amazon Simple Email Service (SES)?**

From the Amazon SES Console, you can set up your receipt rule to have Amazon SES deliver your messages to an AWS Lambda function. The same functionality is available through the AWS SDK and CLI.

## **Q: How do I use an AWS Lambda function to respond to Amazon CloudWatch alarms?**

First, configure the alarm to send Amazon SNS notifications. Then from the AWS Lambda console, select a Lambda function and associate it with that Amazon SNS topic. See the [Amazon CloudWatch Developer Guide](#) for more on setting up Amazon CloudWatch alarms.

## **Q: How do I use an AWS Lambda function to respond to changes in user or device data managed by Amazon Cognito?**

From the AWS Lambda console, you can select a function to trigger when any datasets associated with an [Amazon Cognito](#) identity pool are synchronized. This same functionality is also available through the AWS SDK and CLI. Visit Amazon Cognito for more information on using Amazon Cognito to share and synchronize data across a user's devices.

## **Q: How can my application trigger an AWS Lambda function directly?**

You can invoke a Lambda function using a custom event through AWS Lambda's invoke API. Only the function's owner or another AWS account that the owner has granted permission can invoke the function. Visit the [Lambda Developers Guide](#) to learn more.

## **Q: What is the latency of invoking an AWS Lambda function in response to an event?**

AWS Lambda is designed to process events within milliseconds. Latency will be higher immediately after a Lambda function is created, updated, or if it has not been used recently.

## **Q: How do I create a mobile back-end using AWS Lambda?**

You upload the code you want AWS Lambda to execute and then invoke it from your mobile app using the AWS Lambda SDK included in the AWS Mobile SDK. You can make both direct (synchronous) calls to retrieve or check data in real time as well as asynchronous calls. You can also define a custom API using Amazon API Gateway and invoke your Lambda functions through

any REST compatible client. To learn more about the AWS Mobile SDK, visit the [AWS Mobile SDK](#) page. To learn more about Amazon API Gateway, visit the [Amazon API Gateway](#) page.

### **Q: How do I invoke an AWS Lambda function over HTTPS?**

You can invoke a Lambda function over HTTPS by defining a custom RESTful API using Amazon API Gateway. This gives you an endpoint for your function which can respond to REST calls like GET, PUT and POST. Read more about using AWS Lambda with Amazon API Gateway.

### **Q: How can my AWS Lambda function customize its behavior to the device and app making the request?**

When called through the AWS Mobile SDK, AWS Lambda functions automatically gain insight into the device and application that made the call through the 'context' object.

### **Q: How can my AWS Lambda function personalize their behavior based on the identity of the end user of an application?**

When your app uses the Amazon Cognito identity, end users can authenticate themselves using a variety of public login providers such as Amazon, Facebook, Google, and other OpenID Connect-compatible services. User identity is then automatically and securely presented to your Lambda function in the form of an Amazon Cognito id, allowing it to access user data from Amazon Cognito, or as a key to store and retrieve data in Amazon DynamoDB or other web services.

### **Q: How do I create an Alexa skill using AWS Lambda?**

AWS Lambda is integrated with the Alexa Skills Kit, a collection of self-service APIs, tools, documentation and code samples that make it easy for you to create voice-driven capabilities (or "skills") for Alexa. You simply upload the Lambda function code for the new Alexa skill you are creating, and AWS Lambda does the rest, executing the code in response to Alexa voice interactions and automatically managing the compute resources on your behalf. Read the Alexa Skills Kit documentation for more details.

### **Q: What happens if my function fails while processing an event?**

For Amazon S3 bucket notifications and custom events, AWS Lambda will attempt execution of your function three times in the event of an error condition in your code or if you exceed a service or resource limit.

For ordered event sources that AWS Lambda polls on your behalf, such as Amazon DynamoDB Streams and Amazon Kinesis streams, Lambda will continue attempting execution in the event of a developer code error until the data expires. You can monitor progress through the Amazon

Kinesis and Amazon DynamoDB consoles and through the Amazon CloudWatch metrics that AWS Lambda generates for your function. You can also set Amazon CloudWatch alarms based on error or execution throttling rates.

## Using AWS Lambda to build applications

### Q: What is a serverless application?

Lambda-based applications (also referred to as serverless applications) are composed of functions triggered by events. A typical serverless application consists of one or more functions triggered by events such as object uploads to Amazon S3, Amazon SNS notifications, or API actions. These functions can stand alone or leverage other resources such as DynamoDB tables or Amazon S3 buckets. The most basic serverless application is simply a function.

### Q: How do I deploy and manage a serverless application?

You can deploy and manage your serverless applications using the AWS Serverless Application Model (AWS SAM). AWS SAM is a specification that prescribes the rules for expressing serverless applications on AWS. This specification aligns with the syntax used by AWS CloudFormation today and is supported natively within AWS CloudFormation as a set of resource types (referred to as "serverless resources"). These resources make it easier for AWS customers to use CloudFormation to configure and deploy serverless applications, using existing CloudFormation APIs.

### Q: How can I discover existing serverless applications developed by the AWS community?

You can choose from a collection of serverless applications published by developers, companies, and partners in the AWS community with the [AWS Serverless Application Repository](#). After finding an application, you can configure and deploy it straight from the [Lambda console](#).

### Q: How do I automate deployment for a serverless application?

You can automate your serverless application's release process using AWS CodePipeline and AWS CodeDeploy. CodePipeline is a continuous delivery service that enables you to model, visualize and automate the steps required to release your serverless application. CodeDeploy provides a deployment automation engine for your Lambda-based applications. CodeDeploy lets you orchestrate deployments according to established best-practice methodologies such as canary and linear deployments, and helps you establish the necessary guardrails to verify that newly-deployed code is safe, stable, and ready to be fully released to production.

To learn more about serverless CI/CD, visit our [documentation](#).

## **Q: How do I get started on building a serverless application?**

To get started, visit the AWS Lambda console and download one of our blueprints. The file you download will contain an AWS SAM file (which defines the AWS resources in your application), and a .ZIP file (which includes your function's code). You can then use AWS CloudFormation commands to package and deploy the serverless application that you just downloaded. For more details, visit our [documentation](#).

## **Q: How do I coordinate calls between multiple AWS Lambda functions?**

You can use [AWS Step Functions](#) to coordinate a series of AWS Lambda functions in a specific order. You can invoke multiple Lambda functions sequentially, passing the output of one to the other, and/or in parallel, and Step Functions will maintain state during executions for you.

## **Q: How do I troubleshoot a serverless application?**

You can enable your Lambda function for tracing with [AWS X-Ray](#) by adding X-Ray permissions to your Lambda function's execution role and changing your function's "tracing mode" to "active." When X-Ray is enabled for your Lambda function, AWS Lambda will emit tracing information to X-Ray regarding the Lambda service overhead incurred when invoking your function. This will provide you with insights such as Lambda service overhead, function init time, and function execution time. In addition, you can include the X-Ray SDK in your Lambda deployment package to create your own trace segments, annotate your traces, or view trace segments for downstream calls made from your Lambda function. X-Ray SDKs are currently available for Node.js and Java. Visit [Troubleshooting Lambda-based applications](#) to learn more. AWS X-Ray rates will apply.

## **Q. Can I build serverless applications that connect to relational databases?**

Yes. You can build highly scalable, secure, Lambda-based serverless applications that connect to relational databases using [Amazon RDS Proxy](#), a highly available database proxy that manages thousands of concurrent connections to relational databases. Currently, RDS Proxy supports MySQL and Aurora databases. You can begin using RDS Proxy through the Amazon RDS console or the AWS Lambda console. Serverless applications that use fully managed connection pools from RDS Proxy will be billed according to [RDS Proxy Pricing](#).

## **Q: How is AWS SAM licensed?**

The specification is open sourced under Apache 2.0, which allows you and others to adopt and incorporate AWS SAM into build, deployment, monitoring and management tools with a commercial-friendly license. You can access the AWS SAM repository on GitHub [here](#).

# Container Image Support

## Q: What is Container Image Support for AWS Lambda?

AWS Lambda now enables you to package and deploy functions as container images. Customers can leverage the flexibility and familiarity of container tooling, and the agility and operational simplicity of AWS Lambda to build applications.

## Q: How can I use Container Image Support for AWS Lambda?

You can start with either an AWS provided base images for Lambda or by using one of your preferred community or private enterprise images. Then, simply use Docker CLI to build the image, upload to Amazon ECR, and then create the function by using all familiar Lambda interfaces and tools, such as the AWS Management Console, the AWS CLI, the AWS SDK, AWS SAM and AWS CloudFormation.

## Q: Which container image types are supported?

You can deploy third party Linux base images (e.g. Alpine or Debian) to Lambda in addition to the Lambda provided images. AWS Lambda will support all images based on the following image manifest formats: Docker Image Manifest V2 Schema 2 (used with Docker version 1.10 and newer) or Open Container Initiative (OCI) Spec (v1.0 and up). Lambda supports images with a size of up to 10GB.

## Q: What base images can I use?

AWS Lambda provides a variety of base images customers can extend, and customers can also use their preferred Linux-based images with a size of up to 10GB.

## Q: What container tools can I use to package and deploy functions as container images?

You can use any container tooling as long as it supports one of the following container image manifest formats: Docker Image Manifest V2 Schema 2 (used with Docker version 1.10 and newer) or Open Container Initiative (OCI) Specifications (v1.0 and up). For example, you can use native container tools (i.e. docker run, docker compose, Buildah and Packer) to define your functions as a container image and deploy to Lambda.

## Q: What AWS Lambda features are available to functions deployed as container images?

All existing AWS Lambda features, with the exception of Lambda layers and Code Signing, can be used with functions deployed as container images. Once deployed, AWS Lambda will treat an image as immutable. Customers can use container layers during their build process to include dependencies.

## **Q: Will AWS Lambda patch and update my deployed container image?**

Not at this time. Your image, once deployed to AWS Lambda, will be immutable. The service will not patch or update the image. However, AWS Lambda will publish curated base images for all supported runtimes that are based on the Lambda managed environment. These published images will be patched and updated along with updates to the AWS Lambda managed runtimes. You can pull and use the latest base image from DockerHub or Amazon ECR Public, re-build your container image and deploy to AWS Lambda via Amazon ECR. This allows you to build and test the updated images and runtimes, prior to deploying the image to production.

## **Q: What are the differences between functions created using ZIP archives vs. container images?**

There are three main differences between functions created using ZIP archives vs. container images:

1. Functions created using ZIP archives have a maximum code package size of 250 MB unzipped, and those created using container images have a maximum image size of 10 GB.
2. Lambda uses Amazon ECR as the underlying code storage for functions defined as container images, so a function may not be invocable when the underlying image is deleted from ECR.
3. ZIP functions are automatically patched for the latest runtime security and bug fixes. Functions defined as container images are immutable and customers are responsible for the components packaged in their function. Customers can leverage the AWS provided base images which are regularly updated by AWS for security and bug fixes, using the most recent patches available.

## **Q: Is there a performance difference between functions defined as zip and container images?**

No - AWS Lambda ensures that the performance profiles for functions packaged as container images is the same as for those packaged as ZIP archives, including typically sub-second start up times.

## **Q: How will I be charged for deploying Lambda functions as container images?**

There is no additional charge for packaging and deploying functions as container images to AWS Lambda. When you invoke your function deployed as a container image, you pay the regular price for requests and execution duration. To learn more, visit [AWS Lambda pricing](#). You will be charged for storing your container images in Amazon ECR at the standard ECR prices. To learn more visit [Amazon ECR pricing](#).

## **Q: What is the Lambda Runtime Interface Emulator (RIE)?**

The Lambda Runtime Interface Emulator is a proxy for Lambda's [Runtime API](#) which allows customers to locally test their Lambda function packaged as a container image. It is a lightweight web-server that converts HTTP requests to JSON events and emulates the Lambda Runtime API. It allows you to locally test your functions using familiar tools such as cURL and the Docker CLI (when testing functions packaged as container images). It also simplifies running your application on additional compute services. You can include the Lambda Runtime Interface Emulator in your container image to have it accept HTTP requests natively instead of the JSON events required for deployment to Lambda. This component does not emulate Lambda's orchestrator, or security and authentication configurations. The Runtime Interface Emulator is open sourced on GitHub. You can get started by downloading and installing it on your local machine.

**Q: Why do I need the Lambda Runtime Interface Emulator (RIE) during local testing?**

The Lambda Runtime API in the running Lambda service accepts JSON events and returns responses. The Lambda Runtime Interface Emulator allows the function packaged as a container image to accept HTTP requests during local testing with tools like cURL, and surface them via the same interface locally to the function. It allows you to use the docker run or docker-compose up command to locally test your lambda application.

**Q: What function behaviors can I test locally with the emulator?**

You can use the emulator to test if your function code is compatible with the Lambda environment, runs successfully and provides the expected output. For example, you can mock test events from different event sources. You can also use it to test extensions and agents built into the container image against the Lambda Extensions API.

**Q: How does the Runtime Interface Emulator (RIE) help me run my Lambda compatible image on additional compute services?**

Customers can add the Runtime Interface Emulator as the entry point to the container image or package it as a sidecar to ensure the container image now accepts HTTP requests instead of JSON events. This simplifies the changes required to run their container image on additional compute services. Customers will be responsible for ensuring they follow all security, performance, and concurrency best practices for their chosen environment. RIE is pre-packaged into the AWS Lambda provided images, and is available by default in AWS SAM CLI. Base image providers can use the [documentation](#) to provide the same experience for their base images.

**Q: How can I deploy my existing containerized application to AWS Lambda?**

You can deploy a containerized application to AWS Lambda if it meets the below requirements:

1. The container image must implement the Lambda Runtime API. We have open-sourced a set of software packages, Runtime Interface Clients (RIC), that implement the Lambda Runtime API, allowing you to seamlessly extend your preferred base images to be Lambda compatible.
2. The container image must be able to run on a read-only filesystem. Your function code can access a writable /tmp directory storage of 512 MB. If you are using an image that requires a writable root directory, configure it to write to the /tmp directory.
3. The files required for execution of function code can be read by the default Lambda user. Lambda defines a default Linux user with least-privileged permissions that follows security best practices. You need to verify that your application code does not rely on files that are restricted by other Linux users for execution.
4. It is a Linux based container image.

## Provisioned Concurrency

### Q: What is AWS Lambda Provisioned Concurrency?

Provisioned Concurrency gives you greater control over the performance of your serverless applications. When enabled, Provisioned Concurrency keeps functions initialized and hyper-ready to respond in double-digit milliseconds.

### Q: How do I set up and manage Provisioned Concurrency?

You can configure concurrency on your function through the AWS Management Console, the Lambda API, the AWS CLI, and AWS CloudFormation. The simplest way to benefit from Provisioned Concurrency is by using AWS Auto Scaling. You can use Application Auto Scaling to configure schedules, or have Auto Scaling automatically adjust the level of Provisioned Concurrency in real time as demand changes. To learn more about Provisioned Concurrency, [see the documentation](#).

### Q: Do I need to change my code if I want to use Provisioned Concurrency?

You don't need to make any changes to your code to use Provisioned Concurrency. It works seamlessly with all existing functions and runtimes. There is no change to the invocation and execution model of Lambda when using Provisioned Concurrency.

### Q: How will I be charged for Provisioned Concurrency?

Provisioned Concurrency adds a pricing dimension, of 'Provisioned Concurrency', for keeping functions initialized. When enabled, you pay for the amount of concurrency that you configure

and for the period of time that you configure it. When your function executes while Provisioned Concurrency is configured on it, you also pay for Requests and execution Duration. To learn more about the pricing of Provisioned Concurrency, see [AWS Lambda Pricing](#).

### **Q: When should I use Provisioned Concurrency?**

Provisioned Concurrency is ideal for building latency-sensitive applications, such as web or mobile backends, synchronously invoked APIs, and interactive microservices. You can easily configure the appropriate amount of concurrency based on your application's unique demand. You can increase the amount of concurrency during times of high demand and lower it, or turn it off completely, when demand decreases.

### **Q: What happens if a functions receives invocations above the configured level of Provisioned Concurrency?**

If the concurrency of a function reaches the configured level, subsequent invocations of the function have the latency and scale characteristics of regular Lambda functions. You can restrict your function to only scale up to the configured level. Doing so prevents the function from exceeding the configured level of Provisioned Concurrency. This is a mechanism to prevent undesired variability in your application when demand exceeds the anticipated amount.

## **Amazon EFS for AWS Lambda**

### **Q: What is Amazon EFS for AWS Lambda?**

With Amazon Elastic File System (Amazon EFS) for AWS Lambda, customers can securely read, write and persist large volumes of data at virtually any scale using a fully managed elastic NFS file system that can scale on demand without the need for provisioning or capacity management. Previously, developers added code to their functions to download data from S3 or databases to local temporary storage, limited to 512MB. With EFS for Lambda, developers don't need to write code to download data to temporary storage in order to process it.

### **Q: How do I set up Amazon EFS for Lambda?**

Developers can easily connect an existing EFS file system to a Lambda function via an [EFS Access Point](#) by using the console, CLI or SDK. When the function is first invoked, the file system is automatically mounted and made available to function code. You can learn more in the documentation.

### **Q: Do I need to configure my function with VPC settings before I can use my Amazon EFS file system?**

Yes. Mount targets for Amazon EFS are associated with a subnets in a VPC. The AWS Lambda function needs to be configured to access that VPC.

**Q: Who should use Amazon EFS for Lambda?**

Using EFS for Lambda is ideal for building machine learning applications or loading large reference files or models, processing or backing up large amounts of data, hosting web content, or developing internal build systems. Customers can also use EFS for Lambda to keep state between invocations within a stateful microservice architecture, in a StepFunctions workflow, or sharing files between serverless applications and instance or container-based applications.

**Q: Will my data be encrypted in transit?**

Yes. Data encryption in transit uses industry standard Transport Layer Security (TLS) 1.2 to encrypt data sent between AWS Lambda functions and the Amazon EFS file systems.

**Q: Is my data encrypted at rest?**

Customers can provision Amazon EFS to encrypt data at rest. Data encrypted at rest is transparently encrypted while being written, and transparently decrypted while being read, so you don't have to modify your applications. Encryption keys are managed by the AWS Key Management Service (KMS), eliminating the need to build and maintain a secure key management infrastructure.

**Q: How will I be charged for Amazon EFS for AWS Lambda?**

There is no additional charge for using Amazon EFS for AWS Lambda. Customers pay the standard price for AWS Lambda and for Amazon EFS. When using Lambda and EFS in the same availability zone, customers are not charged for data transfer. However, if they use VPC peering for Cross-Account access, they will incur data transfer charges. To learn more, please see [Pricing](#).

**Q: Can I associate more than one Amazon EFS file system with my AWS Lambda function?**

No. Each Lambda function will be able to access one EFS file system.

**Q: Can I use the same Amazon EFS file system across multiple functions, containers, and instances?**

Yes. Amazon EFS supports Lambda functions, ECS and Fargate containers, and EC2 instances. You can share the same file system and use IAM policy and Access Points to control what each function, container, or instance has access to.

# Lambda Extensions

## Q: What is AWS Lambda Extensions?

AWS Lambda Extensions (Preview) lets you integrate Lambda with your favorite tools for monitoring, observability, security, and governance. Extensions enable you and your preferred tooling vendors to plug into Lambda's lifecycle and integrate more deeply into the Lambda execution environment.

## Q: How do Lambda extensions work?

Extensions are companion processes which run within Lambda's execution environment which is where your function code is executed. In addition, they can run outside of the function invocation - i.e. they start before the function is initialized, run in parallel with the function, can run after the function execution is complete, and can also run before the Lambda service shuts down the execution environment.

## Q: What can I use Lambda extensions for?

You can use extensions for your favorite tools for monitoring, observability, security, and governance from AWS as well as the following partners: AppDynamics, Datadog, Dynatrace, Epsagon, HashiCorp, Lumigo, Check Point CloudGuard, New Relic, Thundra, AWS AppConfig, Amazon CloudWatch. To learn more about these extensions, visit the [launch blog post](#).

## Q: How do I set up and manage Lambda extensions?

You can deploy extensions, using Layers, on one or more Lambda functions using the Console, CLI, or Infrastructure as Code tools such as CloudFormation, the AWS Serverless Application Model, and Terraform. To get started, [visit the documentation](#). During the preview, you can run a maximum of 10 extensions per function.

## Q: What runtimes can I use AWS Lambda extensions with?

You can use extensions with the following runtimes: .NET Core 3.1 (C#/PowerShell) (dotnetcore3.1), Custom runtime (provided), Custom runtime on Amazon Linux 2 (provided.al2), Java 11 (Corretto) (java11), Java 8 (Corretto) (java8.al2), Node.js 12.x (nodejs12.x), Node.js 10.x (nodejs10.x), Python 3.8 (python3.8), Python 3.7 (python3.7), Ruby 2.7 (ruby2.7), Ruby 2.5 (ruby2.5). Lambda Extensions and the functions they're extending can use different runtimes.

## Q: Do Extensions count towards the deployment package limit?

Yes, the total unzipped size of the function and all Extensions cannot exceed the unzipped deployment package size limit of 250 MB.

## **Q: Is there a performance impact of using an extension?**

Extensions may impact the performance of your function because they share resources such as CPU, memory, and storage with the function, and because extensions are initialized before function code. For example, if an extension performs compute intensive operations, you may see your function's execution duration increase because the extension and your function code share the same CPU resources.

You can use the *PostRuntimeExecutionDuration* metric to measure the extra time the extension takes after the function execution, and, you can use the MaxMemoryUsed metric to measure the increase in memory used. To understand the impact of a specific extension, you can also use the Duration metric. Currently, the function execution response is returned after function execution and extension execution have completed. To learn more, visit the [Lambda developer documentation](#).

## **Q: How will I be charged for using Lambda extensions?**

Extensions share the same billing model as Lambda functions. When using Lambda functions with extensions, you pay for requests served and the combined compute time used to run your code and all extensions, in 1ms increments. You will be charged for compute time as per existing Lambda duration pricing. To learn more, see [AWS Lambda pricing](#).

The Lambda lifecycle is made up of three distinct phases: 'init', when AWS Lambda initializes the function, dependencies, and extensions; 'invoke', when Lambda executes function and extension code in response to triggers; and 'shut down', after function execution has completed, but extension code could still be executing. To learn more about the Lambda lifecycle, see the documentation on the Lambda Execution Environment.

*Note: a known issue at the time of this preview announcement is that you are only charged for compute time used during the 'invoke' phase of the Lambda lifecycle. All phases will be billed for on a going forward basis once this issue is corrected; however, you will not be back-billed for any phases that were not billed due to this issue.*

There is no additional cost for installing extensions, although partner offerings may be chargeable. See relevant partner website for details.

## **Q: Can I create my own custom Lambda extensions?**

Yes, by using the AWS Lambda Runtime Extensions API. Visit the documentation to learn more.

## **Q: How do extensions work while Provisioned Concurrency is enabled?**

Provisioned Concurrency keeps functions initialized and ready to respond in double-digit milliseconds. When enabled, Provisioned Concurrency will also initialize extensions and keep them ready to execute alongside function code.

### **Q: What permissions do extensions have?**

Because Extensions are executed within the same environment as a Lambda function. They have access to the same resources as the function and permissions are shared between the function and the extension, therefore they share credentials, role, and environment variables. Extensions have read-only access to function code, and can read and write in /tmp.

### **Q: What is the AWS Lambda Runtime Logs API?**

The AWS Lambda Runtime Logs API enables you to use extensions to send logs from AWS Lambda functions directly to a destination of your choice. Extensions use this API to subscribe to the same logs that are streamed to Amazon CloudWatch Logs, and can then process, filter, and send them to any preferred destination.

### **Q: How does the Runtime Logs API work?**

The Lambda service automatically captures logs and streams them to Amazon CloudWatch. This stream contains the logs which are generated from within your function code, and also those generated by the Lambda service as part of the invoke.

The Runtime Logs API allows extension authors to subscribe to the same log streams directly from within the Lambda execution environment. After receiving the subscription request, the Lambda service streams logs to the extension via HTTP or TCP in addition to sending them to CloudWatch.

### **Q: How do I get started with using the Runtime Logs API?**

You can deploy extensions that use the Runtime Logs API, using Layers, on one or more Lambda functions using the Console, CLI, or Infrastructure as Code tools such as CloudFormation, the AWS Serverless Application Model, and Terraform.

You can use extensions for your favorite tools for monitoring, observability, security, and governance from AWS as well as the following partners: Coralogix, Datadog, Honeycomb, Lumigo, New Relic, Sumo Logic, and Amazon CloudWatch. To learn more about these extensions, visit the [launch blog post](#).

### **Q: Is there a performance impact of using the Runtime Logs API?**

You can only use the Runtime Logs API from within AWS Lambda Extensions. Extensions may impact the performance of your function because they share resources such as CPU, memory,

and storage with the function, and because extensions are initialized before function code. For example, if an extension performs compute intensive operations, you may see your function's execution duration increase because the extension and your function code share the same CPU resources. Additionally, each subscription to the Runtime Logs API could consume additional memory to store logs, on top of what the extension containing it consumes.

#### **Q: How will I be charged for using the Runtime Logs API?**

There is no additional charge for using the AWS Lambda Runtime Logs API. Extensions that make use of the Runtime Logs API share the same billing model as other extensions and Lambda functions. To learn more about Extensions pricing, please [see the FAQs](#).

#### **Q: Does using the Runtime Logs API disable sending logs to Amazon CloudWatch Logs?**

No, by default, the Lambda platform sends all logs to CloudWatch Logs, and using the Runtime Logs API does not disable egress to CloudWatch Logs.

## **Lambda@Edge**

#### **Q: What is Lambda@Edge?**

[Lambda@Edge](#) allows you to run code across AWS locations globally without provisioning or managing servers, responding to end users at the lowest network latency. You just upload your Node.js or Python code to AWS Lambda and configure your function to be triggered in response to [Amazon CloudFront](#) requests (i.e., when a viewer request lands, when a request is forwarded to or received back from the origin, and right before responding back to the end user). The code is then ready to execute across AWS locations globally when a request for content is received, and scales with the volume of CloudFront requests globally. Learn more in our [documentation](#).

#### **Q: How do I use Lambda@Edge?**

To use Lambda@Edge, you just upload your code to AWS Lambda and associate a function version to be triggered in response to Amazon CloudFront requests. Your code must satisfy the Lambda@Edge service limits. Lambda@Edge supports Node.js and Python for global invocation by CloudFront events at this time. Learn more in our [documentation](#).

#### **Q: When should I use Lambda@Edge?**

Lambda@Edge is optimized for latency sensitive use cases where your end viewers are distributed globally. All the information you need to make a decision should be available at the CloudFront edge, within the function and the request. This means that use cases where you are looking to make decisions on how to serve content based on user characteristics (e.g., location,

client device, etc) can now be executed and served close to your users without having to be routed back to a centralized server.

### **Q: Can I deploy my existing Lambda functions for global invocation?**

You can associate existing Lambda functions with CloudFront events for global invocation if the function satisfies the Lambda@Edge service requirements and limits. Read more [here](#) on how to update your function properties.

### **Q: What Amazon CloudFront events can be used to trigger my functions?**

Your functions will automatically trigger in response to the following Amazon CloudFront events:

- Viewer Request - This event occurs when an end user or a device on the Internet makes an HTTP(S) request to CloudFront, and the request arrives at the edge location closest to that user.
- Viewer Response - This event occurs when the CloudFront server at the edge is ready to respond to the end user or the device that made the request.
- Origin Request - This event occurs when the CloudFront edge server does not already have the requested object in its cache, and the viewer request is ready to be sent to your backend origin webserver (e.g. Amazon EC2, or Application Load Balancer, or Amazon S3).
- Origin Response - This event occurs when the CloudFront server at the edge receives a response from your backend origin webserver.

### **Q: How is AWS Lambda@Edge different from using AWS Lambda behind Amazon API Gateway?**

The difference is that API Gateway and Lambda are regional services. Using [Lambda@Edge](#) and [Amazon CloudFront](#) allows you to execute logic across multiple AWS locations based on where your end viewers are located.

## **Scalability and availability**

### **Q: How available are AWS Lambda functions?**

AWS Lambda is designed to use replication and redundancy to provide high availability for both the service itself and for the Lambda functions it operates. There are no maintenance windows or scheduled downtimes for either.

**Q: Do my AWS Lambda functions remain available when I change my code or its configuration?**

Yes. When you update a Lambda function, there will be a brief window of time, typically less than a minute, when requests could be served by either the old or the new version of your function.

**Q: Is there a limit to the number of AWS Lambda functions I can execute at once?**

No. AWS Lambda is designed to run many instances of your functions in parallel. However, AWS Lambda has a default safety throttle for number of concurrent executions per account per region (visit [here](#) for info on default safety throttle limits). You can also control the maximum concurrent executions for individual AWS Lambda functions which you can use to reserve a subset of your account concurrency limit for critical functions, or cap traffic rates to downstream resources.

If you wish to submit a request to increase the throttle limit you can visit our [Support Center](#), click "Open a new case," and file a service limit increase request.

**Q: What happens if my account exceeds the default throttle limit on concurrent executions?**

On exceeding the throttle limit, AWS Lambda functions being invoked synchronously will return a throttling error (429 error code). Lambda functions being invoked asynchronously can absorb reasonable bursts of traffic for approximately 15-30 minutes, after which incoming events will be rejected as throttled. In case the Lambda function is being invoked in response to Amazon S3 events, events rejected by AWS Lambda may be retained and retried by S3 for 24 hours. Events from Amazon Kinesis streams and Amazon DynamoDB streams are retried until the Lambda function succeeds or the data expires. Amazon Kinesis and Amazon DynamoDB Streams retain data for 24 hours.

**Q: Is the default limit applied on a per function level?**

No, the default limit only applies at an account level.

**Q: What happens if my Lambda function fails during processing an event?**

On failure, Lambda functions being invoked synchronously will respond with an exception. Lambda functions being invoked asynchronously are retried at least 3 times. Events from Amazon Kinesis streams and Amazon DynamoDB streams are retried until the Lambda function succeeds or the data expires. Kinesis and DynamoDB Streams retain data for a minimum of 24 hours.

**Q: What happens if my Lambda function invocations exhaust the available policy?**

On exceeding the retry policy for asynchronous invocations, you can configure a “dead letter queue” (DLQ) into which the event will be placed; in the absence of a configured DLQ the event may be rejected. On exceeding the retry policy for stream based invocations, the data would have already expired and therefore rejected.

**Q: What resources can I configure as a dead letter queue for a Lambda function?**

You can configure an Amazon SQS queue or an Amazon SNS topic as your dead letter queue.

## Security and access control

**Q: How do I allow my AWS Lambda function access to other AWS resources?**

You grant permissions to your Lambda function to access other resources using an IAM role. AWS Lambda assumes the role while executing your Lambda function, so you always retain full, secure control of exactly which AWS resources it can use. Visit [Setting up AWS Lambda](#) to learn more about roles.

**Q: How do I control which Amazon S3 buckets can call which AWS Lambda functions?**

When you configure an Amazon S3 bucket to send messages to an AWS Lambda function a resource policy rule will be created that grants access. Visit the [Lambda Developer's Guide](#) to learn more about resource policies and access controls for Lambda functions.

**Q: How do I control which Amazon DynamoDB table or Amazon Kinesis stream an AWS Lambda function can poll?**

Access controls are managed through the Lambda function’s role. The role you assign to your Lambda function also determines which resource(s) AWS Lambda can poll on its behalf. Visit the [Lambda Developer's Guide](#) to learn more.

**Q: How do I control which Amazon SQS queue an AWS Lambda function can poll?**

Access controls can be managed by the Lambda function’s role or a resource policy setting on the queue itself. If both policies are present, the more restrictive of the two permissions will be applied.

**Q: Can I access resources behind Amazon VPC with my AWS Lambda function?**

Yes. You can access resources behind Amazon VPC.

**Q: How do I enable and disable the VPC support for my Lambda function?**

To enable VPC support, you need to specify one or more subnets in a single VPC and a security group as part of your function configuration. To disable VPC support, you need to update the function configuration and specify an empty list for the subnet and security group. You can change these settings using the AWS APIs, CLI, or AWS Lambda Management Console.

#### **Q: Can a single Lambda function have access to multiple VPCs?**

No. Lambda functions provide access only to a single VPC. If multiple subnets are specified, they must all be in the same VPC. You can connect to other VPCs by peering your VPCs.

#### **Q: Can Lambda functions in a VPC also be able to access the internet and AWS Service endpoints?**

Lambda functions configured to access resources in a particular VPC will not have access to the internet as a default configuration. If you need access to external endpoints, you will need to create a [NAT](#) in your VPC to forward this traffic and configure your security group to allow this outbound traffic.

#### **Q: What is Code Signing for AWS Lambda?**

Code Signing for AWS Lambda offers trust and integrity controls which enable you to verify that only unaltered code from approved developers is deployed in your Lambda functions. You can use [AWS Signer](#), a fully-managed code signing service to digitally signed code artifacts and configure your Lambda functions to verify the signatures at deployment. Code Signing for AWS Lambda is currently only available for functions packaged as ZIP archives.

#### **Q: How do I create digitally signed code artifacts?**

You can create digitally signed code artifacts using a [Signing Profile](#) through the AWS Signer console, the Signer API, SAM CLI or AWS CLI. To learn more, please see [documentation for AWS Signer](#).

#### **Q: How do I configure my Lambda functions to enable code signing?**

You can enable code signing by creating a Code Signing Configuration through the AWS Management Console, the Lambda API, the AWS CLI, AWS CloudFormation, and AWS SAM. Code Signing Configuration helps you specify the approved signing profiles and configure whether to warn or reject deployments if signature checks fail. Code Signing Configurations can be attached to individual Lambda functions to enable the code signing feature. Such functions now start verifying signatures at deployment.

#### **Q: What signature checks does AWS Lambda perform on deployment?**

AWS Lambda can perform following signature checks at deployment:

- Corrupt signature - This occurs if the code artifact has been altered since signing.
- Mismatched signature - This occurs if the code artifact is signed by a signing profile that is not approved.
- Expired signature - This occurs if the signature is past the configured expiry date.
- Revoked signature - This occurs if the signing profile owner revokes the signing jobs.

To learn more, please see the [AWS Lambda documentation](#).

#### **Q: Can I enable code signing for existing functions?**

Yes, you can enable code signing for existing functions by attaching a code signing configuration to the function. You can do this using the AWS Lambda console, the Lambda API, the AWS CLI, AWS CloudFormation, and AWS SAM.

#### **Q: Is there any additional cost for using Code Signing for AWS Lambda?**

There is no additional cost when using Code Signing for AWS Lambda. You pay the standard price for AWS Lambda. To learn more, please see [Pricing](#).

## **AWS Lambda functions in Java**

#### **Q: How do I compile my AWS Lambda function Java code?**

You can use standard tools like Maven or Gradle to compile your Lambda function. Your build process should mimic the same build process you would use to compile any Java code that depends on the AWS SDK. Run your Java compiler tool on your source files and include the AWS SDK 1.9 or later with transitive dependencies on your classpath. For more details, see our [documentation](#).

#### **Q: What is the JVM environment Lambda uses for execution of my function?**

Lambda provides the Amazon Linux build of openjdk 1.8.

## **AWS Lambda functions in Node.js**

#### **Q: Can I use packages with AWS Lambda?**

Yes. You can use NPM packages as well as custom packages. Learn more [here](#).

#### **Q: Can I execute other programs from within my AWS Lambda function written in Node.js?**

Yes. Lambda's built-in sandbox lets you run batch ("shell") scripts, other language runtimes, utility routines, and executables. Learn more [here](#).

#### **Q: Is it possible to use native modules with AWS Lambda functions written in Node.js?**

Yes. Any statically linked native module can be included in the ZIP file you upload, as well as dynamically linked modules compiled with an rpath pointing to your Lambda function root directory. Learn more [here](#).

#### **Q: Can I execute binaries with AWS Lambda written in Node.js?**

Yes. You can use Node.js' child\_process command to execute a binary that you've included in your function or any executable from Amazon Linux that is visible to your function. Alternatively several NPM packages exist that wrap command line binaries such as node-ffmpeg. Learn more [here](#).

#### **Q: How do I deploy AWS Lambda function code written in Node.js?**

To deploy a Lambda function written in Node.js, simply package your Javascript code and dependent libraries as a ZIP. You can upload the ZIP from your local environment, or specify an Amazon S3 location where the ZIP file is located. For more details, see our [documentation](#).

## **AWS Lambda functions in Python**

#### **Q: Can I use Python packages with AWS Lambda?**

Yes. You can use pip to install any Python packages needed.

## **AWS Lambda functions in C#**

#### **Q: How do I package and deploy an AWS Lambda function in C#?**

You can create a C# Lambda function using the Visual Studio IDE by selecting "Publish to AWS Lambda" in the Solution Explorer. Alternatively, you can directly run the "dotnet lambda publish" command from the dotnet CLI which has the [Lambda CLI tools patch] installed, which creates a ZIP of your C# source code along with all NuGet dependencies as well as your own published DLL assemblies, and automatically uploads it to AWS Lambda using the runtime parameter "dotnetcore1.0"

## AWS Lambda functions in PowerShell

### Q: How do I deploy AWS Lambda function code written in PowerShell?

A PowerShell Lambda deployment package is a ZIP file that contains your PowerShell script, PowerShell modules that are required for your PowerShell script, and the assemblies needed to host PowerShell Core. You then use the *AWSLambdaPSCore* PowerShell module that you can install from the [PowerShell Gallery](#) to create your PowerShell Lambda deployment package.

## AWS Lambda functions in Go

### Q: How do I package and deploy an AWS Lambda function in Go?

Upload your Go executable artifact as a ZIP file through the AWS CLI or Lambda console and select the go1.x runtime. With Lambda, you can use Go's native tools to build and package your code. For more details, read our [documentation](#).

## AWS Lambda functions in Ruby

### Q: How do I deploy AWS Lambda function code written in Ruby?

To deploy a Lambda function written in Ruby, package your Ruby code and gems as a ZIP. You can upload the ZIP from your local environment, or specify an Amazon S3 location where the ZIP file is located.

## Other topics

### Q: Which versions of Amazon Linux, Node.js, Python, JDK, .NET Core, SDKs, and additional libraries does AWS Lambda support?

You can view the list of supported versions [here](#).

### Q: Can I change the version of Amazon Linux or any language runtime?

No. AWS Lambda offers a single version of the operating system and managed language runtime to all users of the service. You can [bring your own language runtime](#) to use in Lambda.

### Q: How can I record and audit calls made to the AWS Lambda API?

AWS Lambda is integrated with AWS CloudTrail. AWS CloudTrail can record and deliver log files to your Amazon S3 bucket describing the API usage of your account.

**Q: How do I coordinate calls between multiple Lambda functions?**

You can use Amazon Step Functions to coordinate multiple invoking Lambda functions. You can invoke multiple Lambda functions serially, passing the output of one to the other, or in parallel. See our [documentation](#) for more details.

**Q: Does AWS Lambda support Advanced Vector Extensions 2 (AVX2)?**

Yes, AWS Lambda supports the Advanced Vector Extensions 2 (AVX2) instruction set. To learn more about how to compile your application code to target this instruction set for improved performance, visit the [AWS Lambda developer documentation](#).